



CMSC 611 – Advanced Computer Architecture

Class Project

Due: May 2nd, 2013

Objective

To experience the design issues of advanced computer architectures through the design of an analyzer for a simplified MIPS CPU using high level programming languages. The considered MIPS CPU adopts the *CDC 6600 scoreboard* scheme to dynamically schedule instruction execution and employ caches in order to expedite memory access.

Project Statement

Consider a simplified version of the MIPS instruction set architecture shown below in Table 1 and whose formats are provided at the end of this document.

Table 1: Reduced MIPS instruction set

Instruction Class	Instruction Mnemonic
Data Transfers	LW, SW, L.D, S.D
Arithmetic/ logical	DADD, DADDI, DSUB, DSUBI, AND, ANDI, OR, ORI, ADD.D, MUL.D, DIV.D, SUB.D
Control	J, BEQ, BNE
Special purpose	HLT (to stop fetching new instructions)

You need to develop an architecture simulator for the MIPS computer whose organization is shown in Figure 1. The simulator is to accept as an **input a program in the MIPS assembly using the subset of instructions in Table 1**. The **output of simulator will be a file containing the cycle time at which each instruction completes the various stages, and statistics for cache access**. The detailed specifications of the input and output files will be provided later in this document. The following explains the CPU and Memory system:




Figure 1: Block diagram description of a MIPS Computer

Memory: The MIPS machine is assumed to have an **instruction cache (I-Cache) with an access (hit) time of one cycle per word**. The organization of I-Cache is direct-mapped with 4 blocks and the block size is 4 words. In addition, the machine has a **data cache (D-Cache) with hit time of one cycle per word**. D-Cache is a 2-way set associative with a total of four 4-words blocks. A least recently used block replacement strategy is to be applied for the D-Cache. A **write-back strategy is employed with a write-allocate policy**.

The I-Cache and D-Cache are both connected to main memory using a shared bus. **In the case of a cache miss, if main memory is busy serving the other cache we have to wait for it to be free and then start accessing it**. In other words, **latency of the main memory will be dynamic depending on the time of a**

request and the state of previous requests. The main memory is accessible through one-word-wide bus, and its access time is 3 cycles per word. In case there are an I-Cache miss and a D-Cache miss happens in the same cycle, priority will be given to I-Cache.

CPU: The MIPS computer employs the CDC 6600 scoreboard scheme in order to dynamically schedule instruction execution. A basic organization of a MIPS processor with a scoreboard, 2 FL multipliers, 1 FL Divider, 1 FL Adder and an integer ALU, is shown in Figure 2. There are four stages in the scoreboard-based pipeline: Issue, Read operands, Execution, and Write results, which replace the traditional ID, EX and WB stages in MIPS. The scoreboard mechanism can effectively achieve instruction-level parallelism with in-order issue, out-of-order execution and out-of-order completion. All instructions pass through the issue stage in order (in-order issue); however, they can get stalled or bypass each other in the second stage (read operands) and thus enter execution out of order. When a fetched instruction cannot be issued, the next instruction cannot be fetched.

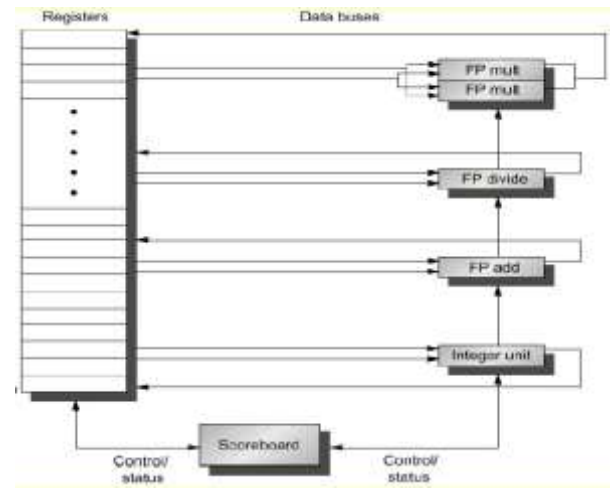


Figure 2: The basic organization of a MIPS processor with a scoreboard

Unconditional jumps complete in the “Issue” stage. The fetched instruction (the next instruction in the program) will be flushed from the “Fetch” stage in that case. In other words a “J” instruction will waste one cycle. On the other hand, conditional branches are resolved in the “Read” stage. Meanwhile the CPU will go ahead and fetch the next instruction, in other words, always “not-taken prediction” will be used in the “Fetch” stage. When a “BEQ” or “BNE” instruction is issued, the pipeline will be stalled until the condition is resolved. In other words, if the branch is “non-taken” one cycle will be wasted. Meanwhile, if we find out in the “Read” stage that the branch is taken, the control unit will flush the “Fetch” stage and update the program counter so the CPU in next cycle will fetch from the branch target address. In other words, if the branch is “Taken”, 2 cycles will be wasted. The branching instructions do not stall because of structural hazard related to the integer unit.

The table below shows the number of cycles each instruction takes in the “Execute” stage.

Instructions	Number of Cycles in “Execute” Stage
HLT, J	0 Cycles (finish in issue stage)
BEQ, BNE	0 Cycle (finish in Read stage)
DADD, DADDI, DSUB, DSUBI, AND, ANDI, OR, ORI	1 Cycle
LW, SW, L.D, S.D	1 Cycle + memory access time (D-Cache)
ADD.D, SUB.D	2 Cycles
MUL.D	30 Cycles
DIV.D	50 Cycles

The number of functional units is to be specified in one of the input files with the name as *config.txt*. The simulator accepts three additional input files; one is for the program containing a mix of the assembly language instructions in Table 1, a file for the initial contents of data memory and another file for the values stored in the integer registers. The output of the simulator should contain the following information:

- The structural and data hazards (RAW, WAR and WAW) in the assembly code which result in pipeline stalls.

- (ii). **The execution time a program takes** (by reporting the cycle number that each instruction completes each stage it passes through).
- (iii). **The performance of the data and instruction caches.** The evaluation criteria should be the total number of cache access requests and the number of the cache hits for the particular cache.

You need to take into account the following additional facts:

- 1) In a MIPS processor with scoreboard, there is **no forwarding hardware**.
- 2) In addition to the **32 word-size registers** (for integers), there are **32 FP registers**; each has 64 bits.
- 3) **Floating point calculations will have no impact on the required output of your simulator.** In fact, you do **not even need to allocate storage** in your simulator for them.
- 4) **Instructions and data are stored in memory starting at address 0x0 and 0x100 respectively.** Load and store instructions **uses word addresses when accessing data.**
- 5) Both conditional and unconditional **jump instructions can be forward and backward.** You can assume that a program will not create a closed loop.
- 6) All caches are blocking and do not support **“hit under misses”**.

The simulator is to be developed in the programming language of your choice. However, you **MUST** submit a “MAKEFILE” that automates the compilation of your project on the *GL machine*. For those using Java or Python, preparing a “MAKEFILE” could be burdensome. In that case, you **MUST** submit a simple shell script file named “make.sh” to automate the compilation. Please also include execution syntax in README file, e.g., “java simulator inputFile.asm data.txt output.txt”. If you use ant and have a build.xml file, please make sure to include it in your project.

Your program must accept the input file in the format specified in the “simulator interface” section. The format of the output must **fully** comply with the specifications.

Simulator Interface

☺ **Input files:** There should be *FOUR* separate input files:

- The first input file is the instruction file “*inst.txt*” that contains the assembly language code, represented as a sequence of instructions in symbolic format such as *L.D* or *ADD.D*, based on the subset of MIPS instructions specified in Table 1 above. **Instructions should be loaded into memory beginning at address 0x00.** Your **simulator should ignore multiple white-spaces and use “,” as the separator for operands.** Moreover, there may be LABELS before some certain instruction so that branch instructions can easily specify the destination. **Every LABEL will be followed by “:” as delimiter.**
- The second input file contains a variable number of 32-bit data words, one per line. These **data words are to be placed in memory beginning at memory location 0x100.** You can assume that the **size of the data segment is 32 words;** meaning that the test cases will not require access to more than 32 words of memory.
- The third input file specifies the initial state of the processor's internal registers. There should be a **sequence of 32 32-bit numbers to be assigned to the integer registers.** These values capture the machine state at the time the assembly code is executed. **The first data word in the file should be loaded into register R₀, the next into register R₁, and so on.** There is no need to preload (or even track the contents of F₀, ..., F₃₁).
- **The fourth input file states the configuration of the scoreboard,** which includes the number of

functional units in the scoreboard. There should be three numbers, one per line, corresponding to the number of units for *FP adder*, *FP Multiplier*, and *FP divider*, respectively. Only one *Integer* unit is allowed in any configuration (and thus there is no need to specify it in the “confi.txt” file).

- All of the four input files and one output file should be specified in the format of *COMMAND-LINE ARGUMENTS*. For instance, a command-line input that starts the simulation program should look like:

simulator inst.txt data.txt reg.txt config.txt result.txt

- The simulator should print the following statement and exit whenever the number of command-line arguments supplied differs from *FOUR*. For example:

```
linux2[1]% ./simulator
Usage: simulator inst.txt data.txt reg.txt config.txt result.txt
```

Note: the role of input files is defined based on their position in argument list. The names and their paths might be different and your simulator should not be restricted to specific name(s) or path(s).

- General comments about the input files format:

- We are not going to test your parser by feeding it with bad input files. However, your program should point out which line of input file is inconsistent with the expected format, and generate an error message and terminate the execution.
- Number of White Spaces (space, tab, enter) should not be a problem for your input parser.
- Your program should not be case sensitive. (e.g., DADD, dadd, dAdd, DAdd, ... are the same)
- You should strongly stick to the format of MIPS instructions. For example if you implement the load immediate as, “LI 1, R1”, it will be wrong (the correct format is LI R1, 1). The table at the end of this document specifies the format of all MIPS instructions covered in this project.

☺ **Output file format:** The simulator must output all desired information to *ONE* output file “*result.txt*”. Be sure to follow the output format *EXACTLY*; deviations from this format will hamper the grading process of your submission. The output should contain the following information:

- The instruction and the clock cycle in which it has completed every stage. If an instruction does not enter a particular stage, leave the entry empty. For example, the “BNE” instruction finishes in the “Read” stage and thus there will be no entries in the following stages for this instruction.
- The existence of RAW, WAR and WAW data hazards that cause STALLS. If an instruction is stalled because such a hazard, mark a “Y” in the corresponding entry, otherwise mark an “N”.
- The existence of structural hazard. If an instruction suffers stalls due to structural hazard, mark a “Y” in the corresponding entry, otherwise mark an “N”.
- The total number of cache access requests for both instruction and data caches.
- The number of cache hits for both instruction and data caches.

- ☒ Use the example below as your guideline for output file format. Do not put extra information such as your name in your output file.

Example:

The following are sample scenarios for which the input files and the expected output are shown.

```
GG:  LD F1, 4(R4)
      LD F2, 8(R5)
      ADD.D F4, F6, F2
      SUB.D F5, F7, F1
      MUL.D F6, F1, F5
      ADD.D F7, F2, F6
      ADD.D F6, F1, F7
      DADDI R4, R4, 2
      DADDI R5, R5, 2
      DSUB R1, R1, R2
      BNE R1, R3, GG
      HLT
```

$$\begin{matrix} 2 \\ 2 \\ 1 \end{matrix}$$
[illegible]

00000000001000000110011000000011
00000000000000000000000000000011
00000000000000000000000000000001
00000000000000000000000000000001
000000000000000000000000100001010
00000000000000000000000010001001
000000000000011101111111010101
00000000000000000001111100100111
0000000000000101010101110101011
000000000000000000010101000000010
00000000000101011111101010101010
00000101011011110101000000000101
01101010101010101011010010110101
00000010101011110001101010101010
00000110101011000010101010101010
00000010101111011010101010011111
00000000000001010101010101101011
00000000000000000000111010101110
000000000000000000000000010101001
00000000000000001110101010101010
01010101011111010101010101011011
00000000000101010101110111010101
00000000111010101010101010101011
00000000000001110101010101010101
00000011010101010100000001010101
000000000000000001010101010010011
001100100100101110101010101001000
01101010101010101011101010101011
00000000101010101010111111101111
0000000000000000000000010110101010
000000000010101010101010101010111
000000000000000000000000000000000

result.txt

Instruction	Fetch	Issue	Read	Exec	Write	RAW	WAR	WAW	Struct
GG: LD F1, 4(R4)	12	13	14	28	29	N	N	N	N
LD F2, 8(R5)	13	30	31	56	57	N	N	N	Y
ADD.D F4, F6, F2	30	31	58	60	61	Y	N	N	N
SUB.D F5, F7, F1	31	32	33	35	36	N	N	N	N
MUL.D F6, F1, F5	43	44	45	75	76	N	N	N	N
ADD.D F7, F2, F6	44	45	77	79	80	Y	N	N	N
ADD.D F6, F1, F7	45	77	81	83	84	Y	N	Y	Y
DADDI R4, R4, 2	77	78	79	80	81	N	N	N	N
DADDI R5, R5, 2	89	90	91	92	93	N	N	N	N
DSUB R1, R1, R2	90	94	95	96	97	N	N	N	Y
BNE R1, R3, GG	94	95	98			Y	N	N	N
HLT	95	96							
GG: LD F1, 4(R4)	99	100	101	115	116	N	N	N	N
LD F2, 8(R5)	100	117	118	132	133	N	N	N	Y
ADD.D F4, F6, F2	117	118	134	136	137	Y	N	N	N
SUB.D F5, F7, F1	118	119	120	122	123	N	N	N	N
MUL.D F6, F1, F5	119	120	124	154	155	Y	N	N	N
ADD.D F7, F2, F6	120	124	156	158	159	Y	N	N	Y
ADD.D F6, F1, F7	124	156	160	162	163	Y	N	Y	N
DADDI R4, R4, 2	156	157	158	159	160	N	N	N	N
DADDI R5, R5, 2	157	161	162	163	164	N	N	N	Y
DSUB R1, R1, R2	161	165	166	167	168	N	N	N	Y
BNE R1, R3, GG	165	166	169			Y	N	N	N
HLT	166	167				N	N	N	N

Total number of access requests for instruction cache: 24

Number of instruction cache hits: 21

Total number of access requests for data cache: 8

Number of data cache hits: 4

Submission Procedure

You can decide on the number of files you would like to submit for this project. However, please make sure that you also provide a **MAKEFILE** for the TA to compile and test your code. For instance, suppose you have just one source code file named as *project.c*, then please write your **MAKEFILE** as the following format, and make sure you provide the *MAKE CLEAN* function:

```
# CMSC 611, Spring 2013, Term Project Makefile

simulator:
    gcc project.c -o simulator

clean:
    -rm simulator *.o core*
```

First you need to ensure the **MAKEFILE** and the source code files are in the same directory. Then run **make**. An executable named *simulator* should appear in the same directory. Ensure that *simulator* runs correctly, and then run **make clean**. Check to ensure that the file *simulator* was deleted from the directory.

Submit *all of your project files* using the GL submission system:

To view the project group for CMSC 611, use the command:

```
> submitproj cs611
```

To submit your project, use the command:

```
> submit cs611 proj
```

To view the files you have submitted, use the command:

```
> submitls cs611 proj
```

To remove the file you already submitted into the system, use the command:

```
> submitrm cs611 proj <file name>
```

For more information for project submission, view <http://www.gl.umbc.edu/submit/>. Please do not email your project submissions to the TA or the instructor, unless the submit command does not work for your account.

Instruction Format and Semantics:

Example Instruction	Instruction Name	Meaning
LW R1, 30(R2)	Load word to an integer register	Regs[R1] <- Mem[30+Regs[R2]]
SW R3, 500(R4)	Store word from an integer register	Mem[500+regs[R4]] <- Regs[R3]
L.D F1, 30(R2)	Load double word to floating point register	Regs[F1] <- Mem[30+Regs[R2]]
S.D F3, 500(R4)	Store double word	Mem[500+regs[R4]] <- Regs[F3]
DADD R1,R2,R3	Add signed	Regs[R1] <- Regs[R2] + Regs[R3]
DADDI R1,R2, 43	Add immediate signed	Regs[R1] <- Regs[R2] + 43
DSUB R1,R2,R3	Subtract signed	Regs[R1] <- Regs[R2] - Regs[R3]
DSUBI R1,R2, 43	Subtract immediate signed	Regs[R1] <- Regs[R2] - 43
AND R1,R2,R3	Bitwise AND	Regs[R1] <- Regs[R2] & Regs[R3]
ANDI R1,R2, 43	Bitwise AND-immediate	Regs[R1] <- Regs[R2] & 3
OR R1,R2,R3	Bitwise OR	Regs[R1] <- Regs[R2] Regs[R3]
ORI R1,R2, 43	Bitwise OR-immediate	Regs[R1] <- Regs[R2] 43
ADD.D F1,F2,F3	Add double word	Regs[F1] <- Regs[F2] + Regs[F3]
SUB.D F1,F2,F3	Subtract double word	Regs[F1] <- Regs[F2] - Regs[F3]
MUL.D F1, F2, F3	Multiply double word	Regs[F1] <- Regs[F2] * Regs[F3]
DIV.D F1, F2, F3	Divide double word	Regs[F1] <- Regs[F2] / Regs[F3]
J LABEL	Unconditional jump	PC <- LABEL
BNE R3, R4, Label	Branch not equal	If(R3 != R4) PC <- Label
BEQ R3, R4, Label	Branch equal	If(R3 == R4) PC <- Label
HLT	Stop fetching new instructions	